

Reinforcement Learning

Kaiwen Zhou

Contents

1 Topic: Markov Decision Processes	1
1.1 Introduction	1
1.2 Goals and Rewards	2
1.3 Episodes	2
1.4 Value Functions	3
1.5 Optimal Policies	4
1.6 The policy improvement theorem	5
1.7 Value Iteration	5
1.8 Generalized policy iteration	6
1.9 Monte Carlo Methods	6
1.10 Policy Gradient Methods	8
1.11 On-policy value function estimation	8
1.12 Stochastic-gradient and Semi-gradient Methods	9
1.13 Semi-gradient control	10
1.14 Policy gradient	10

1 Topic: Markov Decision Processes

MDPs are a classical formalization of sequential decision making, where actions influence not just immediate rewards, but also subsequent situations, or states, and through those future rewards. Thus MDPs involve delayed reward and the need to trade off immediate and delayed reward.

1.1 Introduction

MDPs are a classical formalization of sequential decision making, where actions influence not just immediate rewards, but also subsequent situations, or states, and through those future rewards. Thus MDPs involve delayed reward and the need to trade off immediate and delayed reward. A Markov decision process contains, in particular, a Markov process for a variable that is called the state. If you have ever encountered the concept of a Markov chain described by a transition probability matrix, then you are already familiar with the concept. Similarly, if you have ever used a Kalman filter, you have worked with continuous Markov processes whose probability transition "matrix" is instead a Gaussian density. MDPs are meant to be a straightforward framing of the problem of learning from interaction to achieve a goal. The learner and decision maker is called the agent. The thing it interacts with, comprising everything outside the agent, is called the environment. These interact at each of a sequence of discrete time steps $t = 0, 1, 2, \dots$, the agent selecting actions and the environment responding to these actions and presenting new situations to the agent. The environment also gives rise to rewards, special numerical values that the agent seeks to maximize over time through its choice of actions.




Figure 1: Markov Decision Processes

At each time step t , the agent receives some representation of the environment's state, $S_t \in \mathcal{S}$, and on that basis selects an action, $A_t \in \mathcal{A}(s)$. One time step later, in part as a consequence of its action, the agent receives a numerical reward,

$$R_{t+1} \in \mathcal{R} \subset \mathbb{R}$$

and finds itself in a new state, S_{t+1} .

The way that the agent selects actions is called a policy. Although it is not a particularly clever one, we want to allow as one valid policy, that the agent will roll some combination of dice and/or flip various coins to arrive at the chosen action. In other words a policy may explicitly involve randomness in the choice. Formally, a policy is a mapping from states to a vector of probabilities of selecting each possible action. If the agent is following policy π at time t , then

$$\pi(a \mid s)$$

is the probability that $A_t = a$, given that $S_t = s$.

Any randomness influencing the action choice is independent of the randomness influencing the state-transition dynamics. In a finite MDP, the sets of states, actions, and rewards

$$S, \mathcal{A}, \text{ and } \mathcal{R}$$

all have a finite number of elements.

In this case, the random variables R_t and S_t have well defined discrete probability distributions

$$p(s', r | s, a)$$

dependent only on the preceding state and action:

$$p(s', r | s, a) = \text{prob}(S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a.)$$

The probability function p defines the dynamics of the MDP. In a Markov decision process, the probability of each possible value for the state and the reward,

$$S_t \text{ and } R_t,$$

depends on the immediately preceding state and action,

$$S_{t-1} \text{ and } A_{t-1},$$

and, given the preceding state and action, it depends not at all on earlier states and actions. The state must include information about all aspects of the past agent-environment interaction that make a difference for the future. From the four-argument dynamics function, p , one can (in principle, if there are not too many states) compute anything else one might want to know about the environment, such as the state-transition probabilities

$$p(s' | s, a) = \sum_{r \in \mathcal{R}} p(s', r | s, a)$$

We can also compute the expected rewards for state-action pairs as

$$\mathbb{E}[R_t | S_{t-1} = s, A_{t-1} = a] = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r | s, a)$$

and the expected rewards for state-action-nextstate triples, etc.

1.2 Goals and Rewards

The use of a reward signal to formalize the idea of a goal is one of the most distinctive features of reinforcement learning. Although formulating goals in terms of reward signals might at first appear limiting, in practice it has proved to be flexible and widely applicable. In making a robot learn how to escape from a maze, the reward is often -1 for every time step that passes prior to escape; this encourages the agent to escape as quickly as possible. It is critical that the rewards we set up truly indicate what we want accomplished. For example, a chess-playing agent should be rewarded only for actually winning, not for achieving subgoals such as taking its opponent's pieces or gaining control of the center of the board. If achieving these subgoals were rewarded, then the agent might find a way to achieve them without achieving the real goal; it might find a way to take the opponent's pieces at the cost of losing the game.

1.3 Episodes

In general, we seek to maximize the expected return, where the return, denoted G_t , is defined as some specific function of the reward sequence. In the simplest case the return is the sum of the rewards:

$$G_t = R_{t+1} + R_{t+2} + \cdots + R_T$$

where T is a final time step.

This approach makes sense in applications in which there is a natural notion of final time step, that is, when the agent-environment interaction breaks naturally into subsequences, which we call episodes. Episodes could be plays of a game, trips through a maze, and many other types of sequential play. Each episode ends in a special state called the terminal state, followed by a reset to a standard starting state or to a sample from a standard distribution of starting states. Even if you think of episodes as ending in different ways, such as winning and losing a game, the next episode begins independently of how the previous one ended. Thus the episodes can all be considered to end in the same terminal state, with different rewards for the different outcomes. The time of termination, T , is a random variable that normally varies from episode to episode. On the other hand, in many cases the agent-environment interaction does not break naturally into identifiable episodes, but goes on continually without limit.

For example, this would be the natural way to formulate an on-going process-control task, or an application to a robot with a long life span. We call these continuing tasks. The formulation (10.1) is problematic for continuing tasks, because there is no natural T .

In continuing tasks we often instead define

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+1+k}, \quad \text{for } 0 < \gamma < 1$$

The discount rate determines the present value of future rewards. If

$$0 < \gamma < 1$$

the infinite sum (10.2) has a finite value as long as the reward sequence

$$\{R_k\}$$

is bounded.

Note the useful identity

$$G_t = R_{t+1} + \gamma G_{t+1}$$

which follows directly from the definition.

In fact, there is a way to consider episodic tasks as a special kind of continuing task, which would allow us to work in a unified mathematical notation and framework. We have defined the return as a sum over a finite number of terms in one case (10.1) and as a sum over an infinite number of terms in the other (10.2). These two can be unified by considering episode termination to be the entering of a special absorbing state that transitions only to itself and that generates only rewards of zero. For example, suppose that on a particular run, we get the reward sequence

$$+1, +1, +1, 0, 0, 0, \dots$$

Summing these, we get the same return whether we sum over the first T rewards (here $T = 3$) or over the full infinite sequence. This remains true even if we introduce discounting. With these conventions, we can use the formulation (10.2):

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+1+k}$$

for either continuing or episodic tasks.

In a problem where all episodes terminate, we can set

$$\gamma = 1$$

without any convergence problems.

1.4 Value Functions

The value function of a state s under a policy π , denoted

$$v_{\pi}(s),$$

is the expected return when starting in s and following policy π thereafter.

For MDPs, we can define v_{π} by:

$$v_{\pi}(s) = \mathbb{E}_{\pi} [G_t \mid S_t = s]$$

where the notation \mathbb{E}_{π} means the expectation value in the probability space associated to always following π .

Convince yourself that if a materially-different policy were selected, then the distribution of the random variable G_t will typically be different.

Note that the value of

$$\mathbb{E}_{\pi} [G_t \mid S_t = s]$$

appears superficially to depend on t , but in actual fact it is independent of t due to stationarity and the Markov property.

Note that the value of the terminal state, if any, is always zero. Similarly, we define $q_{\pi}(s, a)$, as the expected return starting from s , taking the action a , and thereafter following policy π :

$$q_{\pi}(s, a) = \mathbb{E}_{\pi} [G_t \mid S_t = s, A_t = a]$$

Now $v_{\pi}(s)$ is called the state-value function and the function

$$q_{\pi}(s, a)$$

is called the action-value function.

The value functions v_{π} and q_{π} can be estimated from experience. For example, if an agent follows policy π and maintains an average, for each state encountered, of the actual returns that have followed that state, then the average will converge to the state's value as the number of times that state is encountered approaches infinity. If separate averages are kept for each action taken in each state, then these averages will similarly converge to the action values. Such estimation methods are Monte Carlo methods. A fundamental property of value functions used throughout

reinforcement learning and dynamic programming is that they satisfy recursive relationships similar to that which we have already established for the return.

For any policy π and any state s , the following consistency condition holds:

$$\begin{aligned} v_\pi(s) &= \mathbb{E}_\pi [G_t \mid S_t = s] \\ &= \mathbb{E}_\pi [R_{t+1} + \gamma G_{t+1} \mid S_t = s] \text{ by (10.3)} \\ &= \sum_a \pi(a \mid s) \sum_{s', r} p(s', r \mid s, a) (r + \gamma \mathbb{E}_\pi [G_{t+1} \mid S_{t+1} = s']) \\ &= \sum_a \pi(a \mid s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_\pi(s')] \end{aligned}$$

This is the Bellman equation for v_π . Intuitively, the Bellman equation expresses a relationship between the value of a state and the values of its successor states. The value function v_π is the unique solution to its Bellman equation.

1.5 Optimal Policies

Value functions define a partial ordering over policies. A policy π is defined to be better than or equal to a policy π' if its expected return is greater than or equal to that of π' for all states, or in other words, if

$$v_\pi(s) \geq v_{\pi'}(s)$$

for all $s \in \mathcal{S}$.

There is always at least one policy that is better than or equal to all other policies. Any such policy is said to be an optimal policy. Although there may be more than one, the notation π_* can be used to denote any policy that is optimal. **All optimal policies share the same state-value function, called the optimal state-value function**, denoted v_* . By the definition of the partial ordering given above, one can easily show that

$$v_*(s) = \max_\pi v_\pi(s) \quad \text{for all } s \in \mathcal{S}$$

The proof is to assume not, and derive a contradiction (exercise).

Optimal policies also share the same optimal action-value function, denoted q_* , which satisfies the analogous relation:

$$q_*(s, a) = \max_\pi q_\pi(s, a) \quad \text{for all } s \in \mathcal{S}, a \in \mathcal{A}(s)$$

For the state-action pair (s, a) , the function value $q_*(s, a)$ gives the expected return for taking action a in state s and thereafter following an optimal policy. Therefore, it follows that

$$q_*(s, a) = \mathbb{E} [R_{t+1} + \gamma v_*(S_{t+1}) \mid S_t = s, A_t = a]$$

This is true because any actions driving the further evolution starting from S_{t+1} would be determined by the optimal policy alone, and not dependent on a in any way at all. This is one place that uses the Markov property. Because v_* is the value function for a policy, it must satisfy the self-consistency condition (10.4) which holds for any such value function. Because it is the optimal value function, however, v_* 's consistency condition can be written in a special form without reference to any specific policy. Intuitively, the value of a state under an optimal policy must equal the expected return for the best action from that state:

$$\begin{aligned} v_*(s) &= \max_{a \in \mathcal{A}(s)} q_{\pi_*}(s, a) \\ &= \max_a \mathbb{E}_{\pi_*} [G_t \mid S_t = s, A_t = a] \\ &= \max_a \mathbb{E}_{\pi_*} [R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a] \text{ by (10.3)} \\ &= \max_a \mathbb{E}_{\pi_*} [R_{t+1} + \gamma v_*(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s')] \end{aligned}$$

This gives the Bellman optimality equation for v_* . As before, v_* is the unique solution to its Bellman equation.

For finite MDPs, suppose that there are n states labeled arbitrarily as $s = 1, \dots, n$, the Bellman optimality equation is actually a system of equations, one for each state

$$\begin{aligned} v_*(s) &= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s')] \\ s &= 1, \dots, n \end{aligned}$$

These are n equations in the n unknowns

$$\{v_*(s) : s = 1, \dots, n\}$$

If the dynamics p of the environment are known, then in principle one can solve this system of equations for v_* and a related system for q_* . The equations are nonlinear due to the presence of the maximum over actions, but numerical solutions are still feasible. Later we will discuss iterative methods which are preferable to direct numerical attack on the Bellman equations.

1.6 The policy improvement theorem

One reason for computing the value function for a policy is to help find better policies. Suppose we have determined the value function v_π for an arbitrary deterministic policy π . For some state s we would like to know whether or not we should change the policy to deterministically choose an action $a \neq \pi(s)$. The action-value function provides a criterion that we can use to decide that. Consider selecting a in state s and thereafter following the policy. The value of this way of behaving is $q_\pi(s, a)$.

Theorem 1.1. *Let π and π' be any pair of deterministic policies such that, for all $s \in \mathcal{S}$,*

$$q_\pi(s, \pi'(s)) \geq v_\pi(s)$$

Then the policy π' must be as good as, or better than π . That is, it must obtain greater or equal expected return from all states:

$$v_{\pi'}(s) \geq v_\pi(s)$$

Moreover, if there is strict inequality of (10.9) at any state, then there must be strict inequality of (10.10) at that state.

Proof.

$$\begin{aligned} v_\pi(s) &\leq q_\pi(s, \pi'(s)) \\ &= \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s, A_t = \pi'(s)] \\ &= \mathbb{E}_{\pi'}[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s] \\ &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma q_\pi(S_{t+1}, \pi'(S_{t+1})) \mid S_t = s] \end{aligned}$$

In the latter, we then replace

$$q_\pi(S_{t+1}, \pi'(S_{t+1}))$$

with

$$\mathbb{E}[R_{t+2} + \gamma v_\pi(S_{t+2}) \mid S_{t+1}, A_{t+1} = \pi'(S_{t+1})]$$

After this replacement, we then have

$$v_\pi(s) \leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 v_\pi(S_{t+2}) \mid S_t = s]$$

If we continue applying this trick recursively, we can show that

$$\begin{aligned} v_\pi(s) &\leq \mathbb{E}_{\pi'}[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 R_{t+4} + \dots \mid S_t = s] \\ &= v_{\pi'}(s) \end{aligned}$$

This completes the proof. □

So far we have seen how, given a policy and its value function, we can easily evaluate a change in the policy at a single state. It is a natural extension to consider changes at all states, selecting at each state the action that appears best according to $q_\pi(s, a)$, in other words, to consider the new greedy policy, π' given by

$$\pi'(s) := \operatorname{argmax}_a q_\pi(s, a)$$

By construction, the greedy policy meets the conditions of the policy improvement theorem (4.7), so we know that it is as good as, or better than, the original policy. The process of making a new policy that improves on an original policy, by making it greedy with respect to the value function of the original policy, is called **policy improvement**. Suppose the new greedy policy π' is as good as, but not better than, the old policy π . Then $v_\pi = v_{\pi'}$ and by definition of the greedy policy,

$$\begin{aligned} v_{\pi'}(s) &= \max_a \mathbb{E}[R_{t+1} + \gamma v_\pi(S_{t+1}) \mid S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_{\pi'}(s')] \end{aligned}$$

But this is the same as the Bellman optimality equation, and therefore, $v_\pi = v_{\pi'} = v_*$. Policy improvement thus must give us a strictly better policy except when the original policy is already optimal.

We will not go through the details, but in fact all the ideas of this section extend easily to stochastic policies. In particular, the policy improvement theorem carries through as stated for the stochastic case. Once a policy, π , has been improved using v_π to yield a better policy, π' , we can then compute $v_{\pi'}$ and improve it again to yield an even better π'' . We can thus obtain a sequence of monotonically improving policies and value functions. This way of finding an optimal policy is called **policy iteration**.

1.7 Value Iteration

One drawback to policy iteration is that each of its iterations involves policy evaluation, which may itself be a protracted iterative computation requiring multiple sweeps through the state set. If policy evaluation is done iteratively, then convergence exactly to v_π occurs only in the limit. Must we wait for exact convergence, or can we stop short of that? In fact, the policy evaluation step of policy iteration can be truncated in several ways without losing the convergence guarantees of policy iteration. One important special case is when policy evaluation is stopped after just one sweep (one update of each state). This algorithm is called value iteration. It can be written as a particularly simple update operation that combines the policy improvement and truncated policy evaluation steps:

$$\begin{aligned}
 v_{k+1}(s) &= \max_a \mathbb{E} [R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t = s, A_t = a] \\
 &= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_k(s')] \\
 &\quad \forall s \in \mathcal{S}
 \end{aligned}$$

For arbitrary v_0 , the sequence $\{v_k\}$ can be shown to converge to v_* under the same conditions that guarantee the existence of v_* . Like policy evaluation, value iteration formally requires an infinite number of iterations to converge exactly to v_* . In practice, we stop once the value function changes by only a small amount in a sweep (i.e. use tolerance).

1.8 Generalized policy iteration

Policy iteration consists of two simultaneous, interacting processes, one making the value function consistent with the current policy (policy evaluation), and the other making the policy greedy with respect to the current value function (policy improvement).

- In policy iteration, these two processes alternate, each completing before the other begins, but this is not really necessary.
- In value iteration, for example, only a single iteration of policy evaluation is performed in between each policy improvement.
- In asynchronous DP methods, the evaluation and improvement processes are interleaved at an even finer grain.

In some cases a single state is updated in one process before returning to the other. As long as both processes continue to update all states, the ultimate result is typically the same-convergence to the optimal value function. The term generalized policy iteration (GPI) refers to the general idea of letting policy-evaluation and policy improvement processes interact, independent of the granularity and other details of the two processes. Almost all reinforcement learning methods are well described as GPI. That is, all have identifiable policies and value functions, with the policy always being improved with respect to the value function and the value function always being driven toward the value function for the policy. The evaluation and improvement processes in GPI can be viewed as both competing and cooperating. They compete in the sense that they pull in opposing directions. Making the policy greedy with respect to the value function typically makes the value function incorrect for the changed policy, and making the value function consistent with the policy typically causes that policy no longer to be greedy. In the long run, however, these two processes interact to find a single joint solution: the optimal value function and an optimal policy.

1.9 Monte Carlo Methods

Monte Carlo methods are ways of solving the reinforcement learning problem based on averaging sample returns. To ensure that well-defined returns are available, here we define Monte Carlo methods only for episodic tasks. Only on the completion of an episode are value estimates and policies changed.

Monte Carlo methods require only experience: sample sequences of states, actions, and rewards from actual or simulated interaction with an environment. Learning from actual experience is striking because it requires no prior knowledge of the environment's dynamics, yet can still attain optimal behavior. Learning from simulated experience is also powerful - although a model is required, the model need only generate sample transitions, not the complete probability distributions of all possible transitions that is required for dynamic programming (DP). In many cases, it is easy to generate experience sampled according to the desired probability distributions, but infeasible to obtain the distributions in explicit form.

We begin by considering Monte Carlo methods for learning the state-value function for a given policy. Recall that the value of a state is the expected return – expected cumulative future discounted reward – starting from that state. An obvious way to estimate it from experience, then, is simply to average the returns observed after visits to that state. As more returns are observed, the average should converge to the expected value. In particular, suppose we wish to estimate $v_\pi(s)$, given a set of episodes obtained by following π and passing through s . **Each occurrence of state s in an episode is called a visit to s .** Of course, s may be visited multiple times in the same episode; let us call the first time it is visited in an episode the first visit to s .

The **first-visit MC method** estimates $v_\pi(s)$ as the average of the returns following first visits to s , whereas the **every-visit MC method** averages the returns following all visits to s . These two Monte Carlo (MC) methods are very similar but have slightly different theoretical properties. First-visit MC has been most widely studied, dating back to the 1940s, and is the one we focus on in this chapter. Every-visit MC extends more naturally to function approximation, as we shall discuss later in the course.

Initialize Returns(s) an empty list, for all $s \in \mathcal{S}$.

- Loop forever (for each episode):
 - (1) Generate an episode following π :
 $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$.
 - (2) set $G = 0$.
 - (3) Loop for each step of episode, $t = T - 1, T - 2, \dots, 0$:
 - (4) (1) set $G = \gamma G + R_{t+1}$
 (2) Unless S_t appears in S_0, S_1, \dots, S_{t-1} :
 (1) append G to Returns(S_t)
 (2) set $V(S_t) = \text{average}(\text{Returns}(S_t))$
-

Every-visit MC would be the same except without the check for S_t having occurred earlier in the episode. Both first-visit MC and every-visit MC converge to $v_\pi(s)$ as the number of visits (or first visits) to s goes to infinity. This is easy to see for the case of first-visit MC. In this case each return is an independent, identically distributed estimate of $v_\pi(s)$ with finite variance. By the law of large numbers the sequence of averages of these estimates converges to their expected value. Each average is itself an unbiased estimate, and the standard deviation of its error falls as $1/\sqrt{n}$ where n is the number of returns averaged. Every-visit MC is less straightforward, but its estimates also converge quadratically to $v_\pi(s)$ (Singh and Sutton, 1996). With a model, state values alone are sufficient to determine a policy; one simply looks ahead one step and chooses whichever action leads to the best combination of reward and next state. Without a model, however, state values alone are not sufficient. One must explicitly estimate the value of each action in order for the values to be useful in suggesting a policy. Thus, one of our primary goals for Monte Carlo methods is to estimate q_* .

To achieve this, we first consider the policy evaluation problem for action values which is to estimate $q_\pi(s, a)$ for a known policy π . The Monte Carlo methods for this are essentially the same as just presented for state values, except now we talk about visits to a state-action pair rather than to a state. A state-action pair s, a is said to be visited in an episode if ever the state s is visited and action a is taken in it. These methods converge quadratically, as before, to the true expected values as the number of visits to each state-action pair approaches infinity. Of course many state-action pairs may never be visited. If π is a deterministic policy, then in following π one will observe returns only for one of the actions from each state. With no returns to average, the Monte Carlo estimates of the other actions will not improve with experience. This is the general problem of maintaining exploration.

One way to do this is by specifying that the episodes start in a state-action pair, and that every pair has a nonzero probability of being selected as the start. This guarantees that all state-action pairs will be visited an infinite number of times in the limit of an infinite number of episodes. We call this the assumption of exploring starts. The assumption of exploring starts cannot be relied upon when learning directly from actual interaction with an environment. In that case the starting conditions are unlikely to be so helpful.

The most common alternative approach to assuring that all state-action pairs are encountered is to consider only policies that are stochastic with a nonzero probability of selecting all actions in each state. We are now ready to consider how Monte Carlo estimation can be used in control, that is, to approximate optimal policies. The overall idea is to proceed according to generalized policy iteration (GPI), which we now describe. In GPI one maintains both an approximate policy and an approximate value function. The value function is repeatedly altered to more closely approximate the value function for the current policy, and the policy is repeatedly improved with respect to the current value function. These two kinds of changes work against each other to some extent, as each creates a moving target for the other.

Consider a Monte Carlo version of classical policy iteration. In this method, we perform alternating complete steps of policy evaluation and policy improvement, beginning with an arbitrary policy π_0 and ending with the optimal policy and optimal action-value function.

$$\pi_0 \xrightarrow{E} q_{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} q_{\pi_1} \xrightarrow{I} \dots$$

where \xrightarrow{E} denotes a complete policy evaluation and \xrightarrow{I} denotes a complete policy improvement.

Policy improvement can be done by constructing each π_{k+1} as the greedy policy with respect to q_k . The policy improvement theorem then applies to π_k and π_{k+1} because for all s ,

$$\begin{aligned} q_{\pi_k}(s, \pi_{k+1}(s)) &= q_{\pi_k}\left(s, \operatorname{argmax}_a q_{\pi_k}(s, a)\right) \\ &= \max_a q_{\pi_k}(s, a) \\ &\geq q_{\pi_k}(s, \pi_k(s)) \\ &= v_{\pi_k}(s) \end{aligned}$$

The policy improvement theorem assures us that each π_{k+1} is uniformly better than π_k , or just as good as π_k , in which case they are both optimal policies. This in turn assures us that the overall process converges to the optimal policy and optimal value function.

In this way Monte Carlo methods can be used to find optimal policies given only sample episodes and no other knowledge of the environment's dynamics. Given a state s , let $\mathcal{A}(s)$ denote the set of actions which are allowed to take from state s , which is for now a finite set with cardinality

$$n_s := |\mathcal{A}(s)| < \infty.$$

Given $\epsilon > 0$, a policy is said to be ϵ -soft if

$$\pi(a | s) \geq \frac{\epsilon}{n_s}$$

for all states s , and for all actions $a \in \mathcal{A}(s)$.

As before, let

$$\operatorname{argmax}_{a \in \mathcal{A}(s)} q(s, a)$$

denote the set of actions $a \in \mathcal{A}(s)$ which maximize the function $a \rightarrow q(s, a)$ over the finite set $\mathcal{A}(s)$.

Sometimes, by abuse of notation, we will write

$$A^* = \operatorname{argmax}_{a \in \mathcal{A}(s)} q(s, a)$$

which typically means that, if the set on the right-hand side has more than one element, then A^* denotes any such element and it is irrelevant which element is chosen.

An ϵ -greedy policy with respect to a given action-value function $q(s, a)$ is defined to be any policy which, from state s , assigns probability

$$1 - \epsilon + \frac{\epsilon}{n_s}$$

to one single action

$$A^* \in \operatorname{argmax}_{a \in \mathcal{A}(s)} q(s, a),$$

and assigns probability

$$\frac{\epsilon}{n_s}$$

to each of the remaining actions.

Since there are $n_s - 1$ actions which are not the chosen greedy action A^* , the total probability works out, as it had better do, to 1.

$$(n_s - 1) \frac{\epsilon}{n_s} + 1 - \epsilon + \frac{\epsilon}{n_s} = 1.$$

Epsilon greedy policies are not unique, but the only ambiguity comes from the need to choose one of the equally-good actions from the set

$$\operatorname{argmax}_{a \in \mathcal{A}(s)} q(s, a)$$

in case the latter has cardinality more than 1.

So you will sometimes see the term "the ϵ -greedy policy" as if it is unique, when we really mean that the arbitrary choice among optimal actions does not matter. It is easy to see that any ϵ -greedy policy is also ϵ -soft (exercise).

The following theorem, which we shall not prove here, is a little more subtle (another exercise).

Theorem 1.2. *For any ϵ -soft policy π , any ϵ -greedy policy with respect to q_π is guaranteed to be better than or equal to π .*

We then have our first reasonably-viable monte carlo control algorithm. Given sufficient computational resources, it will typically achieve the best policy among the ϵ -soft policies.

Initialize Returns(s, a) an empty list, for all $s \in \mathcal{S}, a \in \mathcal{A}(s)$. Loop forever (for each episode):

- Loop forever (for each episode):
 - (1) Generate an episode following π :
 $S_0, A_0, R_1, S_1, A_1, R_2, \dots, S_{T-1}, A_{T-1}, R_T$.
 - (2) set $G = 0$.
 - (3) Loop for each step of episode, $t = T - 1, T - 2, \dots, 0$:
 - (4) (1) set $G = \gamma G + R_{t+1}$
 (2) Unless the pair (S_t, A_t) appears in the sequence of pairs $(S_0, A_0), (S_1, A_1), \dots, (S_{t-1}, A_{t-1})$:
 (1) append G to Returns(S_t, A_t)
 (2) set $Q(S_t, A_t) = \text{average}(\text{Returns}(S_t, A_t))$
 (3) set $A^* = \operatorname{argmax}_a Q(S_t, a)$ and

$$\pi(a | S_t) = \begin{cases} 1 - \epsilon + \epsilon/n_s & a = A^* \\ \epsilon/n_s & a \neq A^* \end{cases}$$

1.10 Policy Gradient Methods

Until now, we have been working with finite state space, meaning that the state $s_t \in \mathcal{S}$ is drawn from a set \mathcal{S} of finite cardinality $|\mathcal{S}|$. That assumption had certain convenient implications. For example, given an action a , we could describe the transition probabilities

$$p(s' | s, a)$$

as an $N \times N$ matrix, where $N = |\mathcal{S}|$.

Even more importantly, in the finite case, it is plausible that we visit the same state multiple times. We used this fact in Monte Carlo value function estimation, where we estimated $v_\pi(s)$ by collecting a list of the returns that were achieved from the first visit to s , out to the end of an episode. There is, of course, the analogous notion of a Markov process on a continuous, real vector space such as \mathbb{R}^d . If you have ever used a Kalman filter, then you are familiar with such processes. If the state space is

$$\mathcal{S} = \mathbb{R}^d,$$

then at best we might hope that we have previously visited a nearby state vector, but the chance that we have visited exactly the same state vector in \mathbb{R}^d becomes vanishingly small.

1.11 On-policy value function estimation

Let's define a quantity called \overline{VE} denoting the value function prediction error. It is defined with respect to a state distribution $\mu(s)$ representing how much we care about the error in each state s . By the error in a state s we mean the square of the difference between the approximate value $\hat{v}(s, w)$ and the true value $v_\pi(s)$. Weighting this over the state space by μ , we obtain a natural objective function, the mean square value error,

$$\overline{VE} := \int_{s \in \mathcal{S}} \mu(s) [v_\pi(s) - \hat{v}(s, w)]^2 ds$$

The square root of this measure, the root VE, gives a rough measure of how much the approximate values differ from the true values and is often used in plots. Often $\mu(s)$ is chosen to be the fraction of time spent in s . Under **on-policy training** this is called the **on-policy distribution**. In continuing tasks, the on-policy distribution is the stationary distribution under π .

1.12 Stochastic-gradient and Semi-gradient Methods

In gradient-descent methods, the weight vector is a column vector with a fixed number of real valued components,

$$\mathbf{w} = (w_1, \dots, w_d)$$

and the approximate value function

$$\hat{v}(s, \mathbf{w})$$

is a differentiable function of \mathbf{w} for all $s \in \mathcal{S}$.

We will be updating \mathbf{w} at each of a series of discrete time steps,

$$t = 0, 1, 2, 3, \dots$$

so we will use notation \mathbf{w}_t for the iterates.

First assume that, on each step, we observe a new example S_t and we observe $v_\pi(S_t)$ consisting of a (possibly randomly selected) state and its true value under the policy. Even though we are given the exact values $v_\pi(S_t)$ there is still a difficult problem because our function approximator has limited resources and thus limited resolution. In particular, there is generally no \mathbf{w} that gets all the states, or even all the examples, exactly correct. In addition, we must generalize to all the other states that have not appeared in examples.

Assume that states appear in examples with the same distribution, μ , over which we are trying to minimize the \overline{VE} . A good strategy in this case is to try to minimize error on the observed examples. Stochastic gradient-descent (SGD) methods do this by adjusting the weight vector after each example by a small amount in the direction that would most reduce the error on that example:

$$\begin{aligned} \mathbf{w}_{t+1} &= \mathbf{w}_t - \frac{1}{2} \alpha \nabla \left[(v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t))^2 \right] \\ &= \mathbf{w}_t + \alpha [v_\pi(S_t) - \hat{v}(S_t, \mathbf{w}_t)] \nabla \hat{v}(S_t, \mathbf{w}_t) \end{aligned}$$

where α is a positive step size and all gradients are with respect to \mathbf{w} .

Convergence results for SGD methods assume that α decreases over time. Note that the step from (11.1) to (11.2) relies on the target being independent of \mathbf{w}_t . We turn now to the case in which the target output, denoted

$$U_t \in \mathbb{R}$$

of the t -th training example is not the true value, $v_\pi(S_t)$, but some, possibly random, approximation to it.

This yields the following update:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [U_t - \hat{v}(S_t, \mathbf{w}_t)] \nabla \hat{v}(S_t, \mathbf{w}_t)$$

If U_t is an unbiased estimate, that is, if

$$\mathbb{E}[U_t \mid S_t = s] = v_\pi(s), \text{ for each } t$$

then \mathbf{w}_t is guaranteed to converge to a local optimum assuming that α decreases over time at the proper rate.

For example, suppose the states in the examples are the states generated by interaction (or simulated interaction) with the environment using policy π . Because the true value of a state is the expected value of the return following it, the Monte Carlo target

$$U_t := G_t$$

is by definition an unbiased estimate.

Bootstrapping methods, such as the TD(0) method we are about to discuss, are not in fact instances of true gradient descent. They take into account the effect of changing the weight vector \mathbf{w}_t on the estimate, but ignore its effect on the target. They include only a part of the gradient and, accordingly, we call them semi-gradient methods. A prototypical semi-gradient method is semi-gradient TD(0), which uses

$$U_t := R_{t+1} + \hat{v}(S_{t+1}, \mathbf{w}_t)$$

as its target. Full pseudocode is given below.

Loop forever (for each episode):

(1) Initialize S .

(2) Loop until S is terminal:

(1) Take action $A \sim \pi(\cdot \mid S)$, observe R, S'

(2) Update

$$\mathbf{w} = \mathbf{w} + \alpha [R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})] \nabla \hat{v}(S, \mathbf{w})$$

(3) Set $S = S'$.

One of the most important special cases of function approximation is that in which the approximate function, $\hat{v}(\cdot, \mathbf{w})$, is a linear function of the weight vector.

$$\hat{v}(s, \mathbf{w}) = \mathbf{w}^\top \mathbf{x}(s)$$

where $\mathbf{x}(s)$ is called a feature vector representing state s .

Mathematically \mathbf{x} maps \mathcal{S} to \mathbb{R}^d . It is natural to use SGD updates with linear function approximation. The gradient of the approximate value function with respect to \mathbf{w} in this case is

$$\nabla \hat{v}(s, \mathbf{w}) = \mathbf{x}(s)$$

Thus, in the linear case the SGD update reduces to a particularly simple form:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [U_t - \hat{v}(s, \mathbf{w}_t)] \mathbf{x}(S_t)$$

Because it is so simple, the linear SGD case is one of the most favorable for mathematical analysis. Many of the useful convergence results for learning systems of all kinds are for linear (or simpler) function approximation methods. Reinforcement learning systems must be capable of generalization if they are to be applicable to artificial intelligence or to large engineering applications. To achieve this, any of a broad range of existing methods for supervised-learning function approximation can be used simply by treating each update as a training example.

1.13 Semi-gradient control

The extension of the semi-gradient prediction methods to action values is straightforward. In this case it is the approximate action-value function,

$$\hat{q} \approx q_\pi,$$

that is represented as a parameterized functional form with weight vector \mathbf{w} .

Whereas before we considered random training examples of the form $S_t \mapsto U_t$, now we consider examples of the form

$$(S_t, A_t) \mapsto U_t$$

The update target U_t can be any approximation of $q_\pi(S_t, A_t)$, including the usual values such as the full Monte Carlo return (G_t). For example, the update for the one-step Sarsa method is

$$\begin{aligned} \mathbf{w}_{t+1} = \mathbf{w}_t + \alpha [R_{t+1} + \gamma \hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}_t) \\ - \hat{q}(S_t, A_t, \mathbf{w}_t)] \nabla \hat{q}(S_t, A_t, \mathbf{w}_t) \end{aligned}$$

For a constant policy, this method converges in the same way that TD(0) does. To form control methods, we need to couple such action-value prediction methods with techniques for policy improvement and action selection. For continuous action space, this is still a topic of intensive ongoing research. On the other hand, if the action set is discrete and not too large, then we can use the techniques already developed. That is, we can compute the greedy action

$$A_{t+1}^* = \operatorname{argmax}_a \hat{q}(S_{t+1}, a, \mathbf{w}_t)$$

Policy improvement is then done (in the on-policy case) by changing the estimation policy to a soft approximation of the greedy policy such as the ϵ -greedy policy. Actions are selected according to this same policy.

1.14 Policy gradient

Consider methods that instead learn a parameterized policy that can select actions without consulting a value function. A value function may still be used to learn the policy parameter, but is not required for action selection. Use the notation $\boldsymbol{\theta} \in \mathbb{R}^d$ for the policy's parameter vector. Thus we write

$$\pi(a | s, \boldsymbol{\theta}) = p(A_t = a | S_t = s, \boldsymbol{\theta}_t = \boldsymbol{\theta})$$

for the probability that action a is taken at time t given that the environment is in state s at time t with parameter $\boldsymbol{\theta}$.

If a method uses a learned value function as well, then the value function's weight vector is denoted $\mathbf{w} \in \mathbb{R}^d$. Consider methods for learning the policy parameter based on maximizing some scalar performance measure $J(\boldsymbol{\theta}) \in \mathbb{R}$. Most methods that we will consider follow the strategy of stochastic gradient ascent, which takes the form

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t + \alpha \widehat{\nabla J(\boldsymbol{\theta}_t)}$$

where

$$\widehat{\nabla J(\boldsymbol{\theta}_t)}$$

is a stochastic estimate whose expectation approximates the gradient of the performance measure with respect to its argument.

Methods that follow this general schema we call policy gradient methods, whether or not they also learn an approximate value function. Methods that learn approximations to both policy and value functions are often called actor-critic methods, where 'actor' is a reference to the learned policy, and 'critic' refers to the learned value function, usually a state-value function.

Now we introduce the most common parameterization for discrete action spaces and point out the advantages it offers over action-value methods. If the action space is discrete and not too large, then a common kind of parameterization is to form parameterized numerical preferences

$$h(s, a, \boldsymbol{\theta}) \in \mathbb{R}$$

for each state-action pair (s, a) .

The actions with the highest preferences in each state are given the highest probabilities of being selected, for example, according to an exponential soft-max distribution:

$$\pi(a | s, \boldsymbol{\theta}) := \frac{e^{h(s, a, \boldsymbol{\theta})}}{\sum_{a'} e^{h(s, a', \boldsymbol{\theta})}}$$

We call this kind of policy parameterization soft-max in action preferences. The action preferences themselves might be computed by a deep artificial neural network (ANN), where $\boldsymbol{\theta}$ is the vector of all the connection weights of the network. Or the preferences could simply be linear in features,

$$h(s, a, \boldsymbol{\theta}) = \boldsymbol{\theta} \cdot \mathbf{x}(s, a)$$

where $\mathbf{x}(s, a)$ are basis functions which then span the space of preference functions of this type.

For policy parameterizations using soft-max in action preferences with linear action preferences, the derivative of the log-probability of a policy takes a particularly simple form. Assuming (11.3) one may show (exercise) that

$$\nabla \ln \pi(a | s, \theta) = \mathbf{x}(s, a) - \sum_{a'} \pi(a' | s, \theta) \mathbf{x}(s, a')$$

A second advantage of parameterizing policies according to the soft-max in action preferences is that it enables the selection of actions with arbitrary probabilities. In problems with significant function approximation, the best approximate policy may be stochastic.

For example, in card games with imperfect information the optimal play is often to do two different things with specific probabilities, such as when bluffing in Poker. In addition to the practical advantages of policy parameterization over ϵ -greedy action selection, there is also an important theoretical advantage.

With continuous policy parameterization the action probabilities change smoothly as a function of the learned parameter, whereas in ϵ -greedy selection the action probabilities may change dramatically for an arbitrarily small change in the estimated action values. Presently we treat the episodic case, for which we define the performance measure as the value of the start state of the episode. We can simplify the notation without losing any meaningful generality by assuming that every episode starts in some particular (non-random) state s_0 . Then, in the episodic case we define performance as

$$J(\theta) = v_{\pi_\theta}(s_0)$$

where v_{π_θ} is the true value function for π_θ , the policy determined by θ .

We will assume no discounting ($\gamma = 1$) for the episodic case. With function approximation it may seem challenging to change the policy parameter in a way that ensures improvement. Fortunately, there is an excellent theoretical answer to this challenge in the form of the policy gradient theorem, which provides an analytic expression for the gradient of performance with respect to the policy parameter (which is what we need to approximate for gradient ascent) that does not involve the derivative of the state distribution.

Theorem 1.3. *In the episodic case, $\nabla J(\theta)$ is proportional to*

$$\int_{s \in \mathcal{S}} \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a | s, \theta) ds$$

where π denotes the policy corresponding to parameter vector θ . The distribution μ here is the on-policy distribution under π .

Stochastic gradient ascent requires a way to obtain samples such that the expectation of the sample gradient is proportional to the actual gradient of the performance measure as a function of the parameter. The sample gradients need only be proportional to the gradient because any constant of proportionality can be absorbed into the step size α . The policy gradient theorem gives an exact expression proportional to the gradient; all that is needed is some way of sampling whose expectation equals or approximates this expression. Notice that the right-hand side of the policy gradient theorem is a sum over states weighted by how often the states occur under the target policy π ; if π is followed, then states will be encountered in these proportions. Thus

$$\begin{aligned} \nabla J(\theta) &\propto \int_{s \in \mathcal{S}} \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a | s, \theta) ds \\ &= \mathbb{E}_\pi \left[\sum_a q_\pi(S_t, a) \nabla \pi(a | S_t, \theta) \right] \end{aligned}$$

Note that we could use (11.5) as the basis for a stochastic gradient descent update, as follows:

$$\theta_{t+1} = \theta_t + \alpha \sum_a \hat{q}(S_t, a, w) \nabla \pi(a | S_t, \theta_t)$$

where $\hat{q}(S_t, a, w)$ is some learned approximation to $q_\pi(S_t, a)$.

This is called an all-actions method.

However, instead of exploring all-actions methods, we would like to find a related SGD method which uses the actual actions chosen while running the simulation of interaction with the environment, rather than requiring us to sum over all the actions. We'd like to replace the sum over actions by an expectation under π . Currently, eqn (11.5) involves an appropriate sum over actions, but each term is not weighted by

$$\pi(a | S_t, \theta)$$

as it would need to be in order to replace the sum by an expectation under π .

So we introduce such a weighting, without changing the equality, by multiplying and then dividing by

$$\pi(a | S_t, \theta)$$

Specifically

$$\begin{aligned} \nabla J(\theta) &\propto \mathbb{E}_\pi \left[\sum_a \pi(a | S_t, \theta) q_\pi(S_t, a) \frac{\nabla \pi(a | S_t, \theta)}{\pi(a | S_t, \theta)} \right] \\ &= \mathbb{E}_\pi \left[q_\pi(S_t, A_t) \frac{\nabla \pi(A_t | S_t, \theta)}{\pi(A_t | S_t, \theta)} \right] \\ &= \mathbb{E}_\pi \left[G_t \frac{\nabla \pi(A_t | S_t, \theta)}{\pi(A_t | S_t, \theta)} \right] \end{aligned}$$

The final expression in brackets is a quantity that can be sampled on each time step whose expectation is proportional to the gradient, so it's a good input for stochastic gradient descent. This leads to a classical algorithm called REINFORCE.

The SGD update is then

$$\theta_{t+1} = \theta_t + \alpha G_t \nabla \ln \pi(A_t | S_t, \theta)$$

This update has an intuitive appeal; the increment $\theta_{t+1} - \theta_t$ is proportional to the product of G_t and a vector. That vector is the gradient of the log-probability of taking the action actually taken, or in other words, that vector is the direction in parameter space that most increases the log-probability of repeating the action A_t on future visits to state S_t . The update increases the parameter vector in this direction proportional to the return, and inversely proportional to the action probability.

The former makes sense because it causes the parameter to move in directions that favor actions that yield the highest return. We can then formalize this into an algorithm, called Monte-Carlo Policy-Gradient Control for π_* , or REINFORCE for short.

Note that REINFORCE uses the complete return from time t , which includes all rewards from $t + 1$ until the end of the episode. In this sense REINFORCE is a Monte Carlo algorithm and is well defined only for the episodic case with all updates made in retrospect after the episode is completed.

Loop forever (for each episode):

(1) Generate an episode

$$S_0, A_0, R_1, \dots, R_T$$

following policy

$$\pi(\cdot \mid \cdot, \theta)$$

(2) Loop for each timestep $t = 0, 1, \dots, T$:

(1)

$$G = \sum_{k=t+1}^T R_k$$

(2)

$$\theta = \theta + \alpha G \nabla \ln \pi(A_t \mid S_t, \theta)$$

As a stochastic gradient method, REINFORCE has good theoretical convergence properties. By construction, the expected update over an episode is in the same direction as the performance gradient. This assures an improvement in expected performance for sufficiently small α , and convergence to a local optimum under standard stochastic approximation conditions for decreasing α . However, as a Monte Carlo method REINFORCE may be of high variance and thus produce slow learning. The convergence of stochastic gradient descent has been studied extensively in the stochastic approximation literature (Bottou, 2012). Convergence results typically require learning rates satisfying the conditions

$$\sum_t \alpha_t^2 < \infty \text{ and } \sum_t \alpha_t = \infty$$

The theorem of Robbins and Siegmund (1985) provides a means to establish almost sure convergence of stochastic gradient descent under surprisingly mild conditions, including cases where the loss function is non-smooth.

Let $b(s)$ be any function or random variable, called a "baseline", that is computable from the state s alone. Then note that the policy gradient theorem still holds with $q_\pi(s, a)$ replaced by $q_\pi(s, a) - b(s)$:

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a (q_\pi(s, a) - b(s)) \nabla \pi(a \mid s, \theta)$$

This follows from the earlier policy gradient theorem, because

$$\sum_a b(s) \nabla \pi(a \mid s, \theta) = 0$$

as can be seen by moving the gradient operator outside the sum. Hence, the baseline leaves the expected value of the update unchanged. This is reassuring because it means we are still doing stochastic gradient descent. It is helpful because it can have a large effect on the variance of the stochastic gradient estimate that we are using in place of the true gradient. One natural choice for the baseline is an estimate of the state value,

$$\hat{v}(S_t; \mathbf{w})$$

where \mathbf{w} is a weight vector learned, e.g. by semi-gradient TD algorithms.

Applying this "natural" choice to the REINFORCE, a.k.a. Monte-Carlo Policy-Gradient Control method discussed previously, gives an enhanced version known as "REINFORCE with baseline" which displays faster learning in many toy models considered in the book.

Loop forever (for each episode):

(1) Generate an episode S_0, A_0, R_1, \dots following policy $\pi(\cdot \mid \cdot, \theta)$

(2) Loop for each timestep t of the episode:

(1) $G = \sum_k R_k$ (sum is for the full episode)

(2) $\delta = G - \hat{v}(S_t, \mathbf{w})$

(3) $\mathbf{w} = \mathbf{w} + \alpha^{(w)} \delta \nabla_{\mathbf{w}} \hat{v}(S_t, \mathbf{w})$

(4) $\theta = \theta + \alpha \delta \nabla \ln \pi(A_t \mid S_t, \theta)$

where $\alpha^{(w)}$ is a different learning rate for \mathbf{w} .

Note the similarity of the \mathbf{w} -update to the gradient monte carlo algorithm for estimating state values discussed above, and in Chapter 9 of Sutton and Barto.

References

- [1] Léon Bottou. Stochastic gradient descent tricks. In *Neural networks: Tricks of the trade*. Springer, 2012, pp. 421-436.
- [2] Herbert Robbins and David Siegmund. A convergence theorem for non negative almost supermartingales and some applications. In *Herbert Robbins Selected Papers*. Springer, 1985, pp. 111-135.